

Data Analytics Project: Operation Analytics and Investigating Metric Spike

Project Description:

Operation Analytics is the analysis done for the complete end to end operations of a company. It is used to find the areas of improvement, predict the overall growth or decline of a company's fortune.

Investigating metric spike is an important part of operation analytics to understand the growth, decline and other insights of the data.

This project insights are to be used by the ops team, support team, marketing team.

Approach:

Based on the provided dataset, and the questions asked by different departments, firstly, the tables were studied, to understand the given data and the relations between each table. Multiple queries were run to get the necessary insights from the data.

Tech-Stack Used:

MySQL Workbench 8.0 CE was used to execute this project. As it provided enough useful set of functions. Also, the connectivity, speed and security of MySQL was suitable for accessing given dataset and run the queries. Also, MS Excel was used for the validation of the outputs derived and for graphical representation of output derived.

Process:

Database named dskdb was created for this project. Tables were created by directly importing the csv files in MySQL Workbench into dskdb database, the type of each column was specified before importing the data.

Insights:

Analysis Performed:

A) Case Study 1 (Job Data):

1. Number of jobs reviewed:

- For this task, we have to the number of jobs reviewed per hour per day for November 2020
- Below query was used to get the total number of jobs reviewed every day, and the time spent on hourly basis per day on the review.

Query:

```
use dskdb;
```

```
select distinct ds as "Date",
```

```
count(job_id) as "Total Number of Jobs Reviewed",
```

sum(time_spent)/3600 as "Time_Spent_in_Hours",

round(count(distinct job_id)/(sum(time_spent)/3600)) as "Jobs Reviewed Per Hour Per Day"
from job_data

group by ds;

Output:

| | Date | Total Number of Jobs Reviewed | Time_Spent_in_Hours | Jobs Reviewed Per Hour Per Day |
|---|---------------------|-------------------------------|---------------------|--------------------------------|
| ► | 2020-11-25 00:00:00 | 1 | 0.0125 | 80 |
| | 2020-11-26 00:00:00 | 1 | 0.0156 | 64 |
| | 2020-11-27 00:00:00 | 1 | 0.0289 | 35 |
| | 2020-11-28 00:00:00 | 2 | 0.0092 | 218 |
| | 2020-11-29 00:00:00 | 1 | 0.0056 | 180 |
| | 2020-11-30 00:00:00 | 2 | 0.0111 | 180 |

From the output table, we conclude that on 28 Nov 2020 maximum number of jobs were reviewed that is 218, and on 27 Nov 2020 least number of jobs were reviewed that is 35.

2. Throughput:

It is the no. of events happening per second.

The question is

- To calculate 7 day rolling average of throughput?
- For throughput, do you prefer daily metric or 7-day rolling and why?

For this, calculation we use below query:

Query:

SELECT ds as "Date", (count(event)/sum(time_spent)) as "Daily Average of Throughput",

SUM((count(event)/sum(time_spent))) OVER (ORDER BY ds) AS "Rolling Total of Throughput",

avg((count(event)/sum(time_spent))) over (order by ds rows between 6 preceding and current row) as "7 day Rolling Average" FROM job_data group by ds;

Output:

| | Date | Daily Average of Throughput | Rolling Total of Throughput | 7 day Rolling Average |
|---|---------------------|-----------------------------|-----------------------------|-----------------------|
| ► | 2020-11-25 00:00:00 | 0.0222 | 0.0222 | 0.02222222 |
| | 2020-11-26 00:00:00 | 0.0179 | 0.0401 | 0.02003968 |
| | 2020-11-27 00:00:00 | 0.0096 | 0.0497 | 0.01656492 |
| | 2020-11-28 00:00:00 | 0.0606 | 0.1103 | 0.02757520 |
| | 2020-11-29 00:00:00 | 0.0500 | 0.1603 | 0.03206016 |
| | 2020-11-30 00:00:00 | 0.0500 | 0.2103 | 0.03505013 |

As seen in the output table above, daily average results have huge difference between days, whereas 7 day rolling average shows a stable graph for the throughput. That's why, 7 day rolling average is to be preferred.

3. Percentage share of each language:

The question is to find the percentage share of each language in last 30 days
For this task, following query was used:

Query:

```
use dskdb;
SELECT distinct language, count (*) as "Content Count",
sum (count (*)) over () as "Total Content",
(count (*)/sum (count (*)) over () *100) as "Percentage"
FROM dskdb.job_data
group by language;
```

Output:

| | language | Content Count | Total Content | Percentage |
|---|----------|---------------|---------------|------------|
| ▶ | English | 1 | 8 | 12.5000 |
| | Arabic | 1 | 8 | 12.5000 |
| | Persian | 3 | 8 | 37.5000 |
| | Hindi | 1 | 8 | 12.5000 |
| | French | 1 | 8 | 12.5000 |
| | Italian | 1 | 8 | 12.5000 |

As shown in the table above, Persian language has highest share with 37.5 %, all other language has 12.5% share.

4. Duplicate rows:

For this task, we check for the duplicate rows in the given table by using the query below:

Query:

```
use dskdb;
select * from
(
select *, row_number()over(partition by ds,job_id,actor_id,event,language,time_spent, org) as
rownum from job_data ) a where rownum>1;
;
```

Output:

| ds | job_id | actor_id | event | language | time_spent | org | rownum |
|----|--------|----------|-------|----------|------------|-----|--------|
|----|--------|----------|-------|----------|------------|-----|--------|

The above output table with no records states that there are no duplicate rows in the table.

However, individual columns having unique ids like column job_id and column actor_id, may have duplicates, which can be checked using the query below:

Query (a):




```
use dskdb;
```

```

select * from
(
select *, row_number()over(partition by job_id) as rownum
from job_data
) a
where rownum>1;
;

```

Output:

| Result Grid  Filter Rows: <input type="text"/> Export:  Wrap Cell Content:  | | | | | | | | |
|--|---------------------|--------|----------|----------|----------|------------|-----|--------|
| | ds | job_id | actor_id | event | language | time_spent | org | rownum |
| ▶ | 2020-11-28 00:00:00 | 23 | 1005 | transfer | Persian | 22 | D | 2 |
| | 2020-11-26 00:00:00 | 23 | 1004 | skip | Persian | 56 | A | 3 |

As shown in the output above, job_id 23 has 2 duplicates in the column.




Query (b):

```

select * from
(
select *, row_number()over(partition by actor_id) as rownum
from job_data
) a
where rownum>1;
;

```

Output:

| Result Grid  Filter Rows: <input type="text"/> Export:  Wrap Cell Content:  | | | | | | | | |
|--|---------------------|--------|----------|----------|----------|------------|-----|--------|
| | ds | job_id | actor_id | event | language | time_spent | org | rownum |
| ▶ | 2020-11-25 00:00:00 | 20 | 1003 | transfer | Italian | 45 | C | 2 |

As shown in the table, actor_id 1003 is duplicate in the column

B) Case Study 2 (Investigating metric spike):

1. User Engagement:

For this task, we calculate the weekly user engagement, by using the query given below:

Query:

```
SELECT week(occurred_at) as Week,event_name,  
count(distinct user_id) as "User Engagement"  
FROM dskdb.events  
group by week(occurred_at), event_name;
```

The output:

| | Week | event_name | User Engagement |
|---|------|------------|-----------------|
| ▶ | 30 | login | 1467 |
| | 29 | login | 1376 |
| | 27 | login | 1372 |
| | 28 | login | 1365 |
| | 30 | home_page | 1362 |
| | 26 | login | 1302 |
| | 31 | login | 1299 |
| | 27 | home_page | 1282 |
| | 24 | login | 1275 |
| | 25 | login | 1264 |
| | 28 | home_page | 1259 |
| | 29 | home_page | 1257 |
| | 23 | login | 1232 |
| | 32 | login | 1225 |
| | 33 | login | 1225 |
| | 26 | home_page | 1204 |

The above output table states that in week 30, highest number of users, i.e 1467 users logged in, the table also shows that the highly used service is login and home_page.

2. User Growth:

For this task, we find the amount of users growing over time for a product, using the below query:

Query:

```
SELECT Year, Month, Users,Active_Users,  
sum(Users) over (order by Year,Month rows between unbounded preceding and current row) as  
Total_Users_Growth,  
sum(Active_Users) over (order by Year,Month rows between unbounded preceding and current  
row) as Active_Users_Growth  
from ( select year(created_at) as Year, month(created_at) as Month,  
count(distinct user_id) as Users,
```

```
count(distinct( case when state="active" then user_id end)) as Active_Users

from dskdb.users


group by Year, Month) a;
```

The output:

| | Year | Month | Users | Active_Users | Cummulative_Users_Total | Cummulative_Active_Total | User_Growth | Active_User_Growth |
|---|------|-------|-------|--------------|-------------------------|--------------------------|-------------|--------------------|
| ▶ | 2013 | 1 | 332 | 160 | 332 | 160 | NULL | NULL |
| | 2013 | 2 | 328 | 160 | 660 | 320 | -4 | 0 |
| | 2013 | 3 | 383 | 150 | 1043 | 470 | 55 | -10 |
| | 2013 | 4 | 410 | 181 | 1453 | 651 | 27 | 31 |
| | 2013 | 5 | 486 | 214 | 1939 | 865 | 76 | 33 |
| | 2013 | 6 | 485 | 213 | 2424 | 1078 | -1 | -1 |
| | 2013 | 7 | 608 | 284 | 3032 | 1362 | 123 | 71 |
| | 2013 | 8 | 636 | 316 | 3668 | 1678 | 28 | 32 |
| | 2013 | 9 | 699 | 330 | 4367 | 2008 | 63 | 14 |
| | 2013 | 10 | 826 | 390 | 5193 | 2398 | 127 | 60 |
| | 2013 | 11 | 816 | 399 | 6009 | 2797 | -10 | 9 |
| | 2013 | 12 | 972 | 486 | 6981 | 3283 | 156 | 87 |
| | 2014 | 1 | 1083 | 552 | 8064 | 3835 | 111 | 66 |
| | 2014 | 2 | 1054 | 525 | 9118 | 4360 | -29 | -27 |
| | 2014 | 3 | 1231 | 615 | 10349 | 4975 | 177 | 90 |
| | 2014 | 4 | 1419 | 726 | 11768 | 5701 | 188 | 111 |
| | 2014 | 5 | 1597 | 779 | 13365 | 6480 | 178 | 53 |

The above output shows the growth of all users and active users.

When put in excel graph, it shows a clear growth rate as shown in below graph:



The above graph was derived using Excel charts, to show the growth over the years

3. Weekly Retention:

The task is to calculate the weekly retention of users-sign up cohort. For this first we get the list of users who have completed the sign-up using the query below:

Query:

```
SELECT * FROM dskdb.events  
where event_type="signup_flow" and event_name="complete_signup";
```

The output:

| user_id | occurred_at | event_type | event_name | location | device | user_type |
|---------|---------------------|-------------|-----------------|----------------|-------------------|-----------|
| 11768 | 2014-05-01 08:03:00 | signup_flow | complete_signup | France | macbook pro | 3 |
| 11770 | 2014-05-01 06:08:00 | signup_flow | complete_signup | Japan | iphone 5s | 3 |
| 11775 | 2014-05-01 16:38:00 | signup_flow | complete_signup | United Kingdom | lenovo thinkpad | 2 |
| 11778 | 2014-05-01 18:49:00 | signup_flow | complete_signup | Indonesia | iphone 4s | 3 |
| 11779 | 2014-05-01 18:24:00 | signup_flow | complete_signup | Germany | samsung galaxy s4 | 1 |
| 11780 | 2014-05-01 10:34:00 | signup_flow | complete_signup | United States | iphone 4s | 3 |

The above table gives the list of users who have completed sign-up.

Now to calculate the weekly retention of signed-up users, we use the query below:

Query:

```
select week(occurred_at) as Week, count(distinct user_id) as "Users Weekly Retented" from  
dskdb.events  
where event_type="engagement"  
and user_id in  
(  
SELECT user_id FROM dskdb.events  
where event_type="signup_flow" and event_name="complete_signup"  
)  
group by week(occurred_at)  
;
```

The output:

| Week | Users Weekly Retented |
|------|-----------------------|
| 17 | 72 |
| 18 | 222 |
| 19 | 323 |
| 20 | 407 |
| 21 | 444 |
| 22 | 515 |
| 23 | 541 |
| 24 | 600 |
| 25 | 616 |
| 26 | 644 |
| 27 | 716 |
| 28 | 725 |
| 29 | 743 |
| 30 | 809 |
| 31 | 734 |
| 32 | 763 |

The above table shows the numbers of users getting retained weekly after signing-up for a product.

4. Weekly Engagement:

The question is to find the weekly engagement per device. For this we use the query below:

Query:

```
SELECT week(occurred_at) as Week, device, count(distinct user_id) as "Engagement"  
FROM dskdb.events  
group by Week, device;
```

The output:

| | Week | device | Engagement |
|---|------|------------------------|------------|
| ► | 17 | acer aspire desktop | 12 |
| | 17 | acer aspire notebook | 23 |
| | 17 | amazon fire phone | 4 |
| | 17 | asus chromebook | 23 |
| | 17 | dell inspiron desktop | 20 |
| | 17 | dell inspiron notebook | 48 |
| | 17 | hp pavilion desktop | 17 |
| | 17 | htc one | 19 |
| | 17 | ipad air | 29 |
| | 17 | ipad mini | 19 |
| | 17 | iphone 4s | 27 |
| | 17 | iphone 5 | 69 |
| | 17 | iphone 5s | 47 |
| | 17 | kindle fire | 6 |
| | 17 | lenovo thinkpad | 94 |
| | 17 | mac mini | 7 |
| | 17 | macbook air | 61 |

The above table shows the weekly engagement per device, device macbook has the highest engagement.

5. Email Engagement:

For this task the question is to find the users engaging with email service. Below query is used for that:

Query:

```
SELECT action, week(occurred_at) as week,  
count(distinct user_id) as "Users_Engagement" FROM dskdb.email_events  
group by action, week(occurred_at)  
order by action, week(occurred_at)  
;
```

The output:

| action | week | Users_Engagement |
|--------------------|------|------------------|
| email_clickthrough | 17 | 166 |
| email_clickthrough | 18 | 425 |
| email_clickthrough | 19 | 476 |
| email_clickthrough | 20 | 501 |
| email_clickthrough | 21 | 436 |
| email_clickthrough | 22 | 478 |
| email_clickthrough | 23 | 529 |
| email_clickthrough | 24 | 549 |
| email_clickthrough | 25 | 524 |
| email_clickthrough | 26 | 550 |
| email_clickthrough | 27 | 613 |

| action | week | Users_Engagement |
|--------------------|------|------------------|
| sent_weekly_digest | 34 | 4111 |
| sent_weekly_digest | 33 | 4012 |
| sent_weekly_digest | 32 | 3897 |
| sent_weekly_digest | 31 | 3793 |
| sent_weekly_digest | 30 | 3706 |
| sent_weekly_digest | 29 | 3592 |
| sent_weekly_digest | 28 | 3499 |
| sent_weekly_digest | 27 | 3399 |
| sent_weekly_digest | 26 | 3302 |
| sent_weekly_digest | 25 | 3207 |
| sent_weekly_digest | 24 | 3105 |

The above output shows the email enagement metrics, as shown in the table
sent_weekly_digest has highest enegeement

Insights:

This project is based on 2 datasets. The first dataset is the data for actors with details on the jobs assigned to actors, time spent, language and other details. The outputs derived from this dataset help understand throughput of the events, the share of language and number of jobs reviewed.

The second dataset is the data of an email service used by the users. The tasks and questions answered above help gain the insights on the user behaviour, user engagement with the service and service performance.

The graph derived above on user growth shows the growth rate over the period of 2 years.

Result:

This project has helped me learn how to import huge dataset, clean the dataset and how to get insights from the data. It has helped me practice the queries and get hands on experience on analysing the data as well as how to derive meaningful graphs from the tables.